

TOPIC 1 NEURAL NETWORKS

Neural Networks

- Multi-Level Perceptrons
- Convolutional Neural Networks
- In 2015 Google released an opensource tool for training neural nets called **TensorFlow** for python
- TensorFlow does all the backend work
- **Keras** is a modeling package that makes it easy to formulate neural nets and then passes them to TensorFlow for training
- conda install keras
- Or just do everything in colab

External Resources

- I don't know a great (free) text resource that explains this content very well
- I did find some pretty good youtube videos though
- <https://www.3blue1brown.com/videos>
 - Scroll down to the 'Neural networks' section
 - There is a 4 part video series that explains multi-level perceptrons well
- <https://www.youtube.com/watch?v=FmpDlaiMleA>
 - This does an OK job at explaining convolutional neural networks

Neural Networks


- **Neural Networks** have become increasingly popular in the past couple decades or so
- The application area that has had the most success is **image classification**
 - What address is written on this envelope?
 - Hotdog or Not Hotdog?
- We will work with a data set of 60,000 images (28x28 pixels each) of handwritten numbers 0-9
 - Our goal is to take the image as input and output which number we think it is
 - The images come from the post office
 - MNIST – The “Hello World!” of image classification

Neural Networks

- A neural net is just a fancy tool for non-linear regression and classification
- Much like trees, there's no closed form solution to the parameters, but we can still get close to them
- The tools we use to fit neural networks are
 - Back Propagation
 - Stochastic Gradient Descent
- These are simple tools that are very powerful
- The first type of neural network we'll study is called
 - Multilevel Perceptron Network
 - Fully Connected (dense) Layers

- A neural network is simply a composition of several functions with parameters we estimate from data

- Input Layer
- Neurons
- Hidden Layers
- Weights
- Biases
- Activation Function
- Output Layer



Neural Networks

- Our job is to find all the weights and biases
 - Assume network structure and activation functions are given
- The number of neurons in the first layer is the dimension of the input data (number of regressors: m)
- No limit to the number of hidden layers or neurons
 - More layers and neurons means more parameters
- The number of neurons in the last layer is the dimension of the output data
 - Usually 1 for quantitative variables, like regression
 - Number of categories for classification
 - Output is “probability” of each category

Neural Networks

- There are M input neurons
- Each input neuron has a weight associated with each of the N neurons on the second layer, $v_{n,m}$
- Each of the N neurons on the second layer also have a bias, b_n
- Every neuron on the second layer shares one activation function, $a_2(\cdot)$
- The neurons on layer 2 are equal to
 - $x_n = a_2(b_n + \sum_{m=1}^M v_{n,m} z_m)$

Neural Networks

- There are N neurons on layer 2
- Each layer 2 neuron has a weight associated with each of the K neurons on the third layer, $w_{k,n}$
- Each of the K neurons on the third layer also have a bias, c_k
- Every neuron on the third layer shares one activation function, $a_3(\cdot)$
- The neurons on layer 3 are equal to
 - $y_k = a_3(c_k + \sum_{n=1}^N w_{k,n}x_n)$
- And so on...

Neural Networks

- This is highly non-linear and non-convex because the activation functions are non-linear
- Some popular activation functions are
 - **ReLU**: $a(x) = \max(x, 0)$
 - **Sigmoid**: $a(x) = \frac{1}{1+e^{-x}}$
 - **SoftMax**: $a(x_i) = \frac{e^{x_i}}{\sum_{j=1}^J e^{x_j}}$
 - **SoftPlus**: $a(x) = \log(1 + e^x)$
 - ...

Neural Networks

- To find the weights and biases we must define an objective function (**loss function**)
- The most popular loss function is **mean-squared-error**
 - For classification problems: **one hot encoding**
- Suppose the output layer has J nodes \hat{y}_j
- The data output is y_{ji}
- $$Loss = \frac{1}{n} \sum_{i=1}^n \sum_{j=1}^J (y_{ji} - \hat{y}_{ji})^2$$
- The objective is to minimize the MSE by finding the optimal weights and biases
- There are many other popular loss functions
 - **Categorical Cross Entropy:** $\frac{-1}{n} \sum_{i=1}^n \sum_{j=1}^J y_{ji} \log(p_{ji})$


Back Propagation

- The most common way to do this is **back propagation**
 - This is the NN word for **stochastic gradient descent**
- Start with a random guess of weights and biases
- Calculate the derivative of the loss function with respect to every weight and bias
 - The vector of these derivatives is the **gradient**
- Take a small step in the direction of the -gradient
 - We talked about variations of this last semester
- This relies on being able to calculate the derivatives
 - It turns out this is just the chain rule

Depth Revolution

- In the early 2010's teams participating in the ImageNet competition started using **deep** neural networks
 - They drastically changed the study of NN's

Revolution of Depth



Depth Revolution

- Neural Nets get their name from the brain
- People used to use indicator functions as activation

$$- I(x) = \begin{cases} 1, & x \geq 0 \\ 0, & \text{else} \end{cases}$$


- Training this is hard: the derivative is always zero
- The sigmoid was used as a smooth approximation
 - The derivative is always in (0,1)
- If you have a deep network you have to multiply lots of numbers less than 1:
vanishing derivative
- This is why ReLU is popular – derivative is 1 or 0
- We can now train deep networks!

Neural Networks


- With deep networks the number of parameters to estimate grows very quickly
 - This easily leads to over fitting
- There are a couple ways to avoid this
- We can use a Lasso or Ridge term
 - Penalize our objective with the sum of squared or absolute valued parameters
 - In the NN space these are called **regularizers**
- Each layer can have it's own penalty parameter, λ

Dropout

- A popular alternative in NN's is called **dropout**
 - On a step of SGD just randomly set some neurons equal to zero
 - On another step set a different group of neurons equal to zero




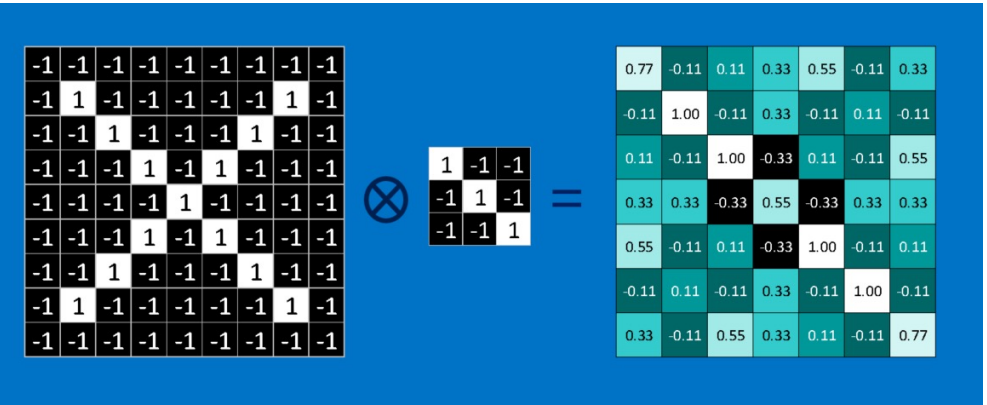


(a) Standard Neural Net



(b) After applying dropout.

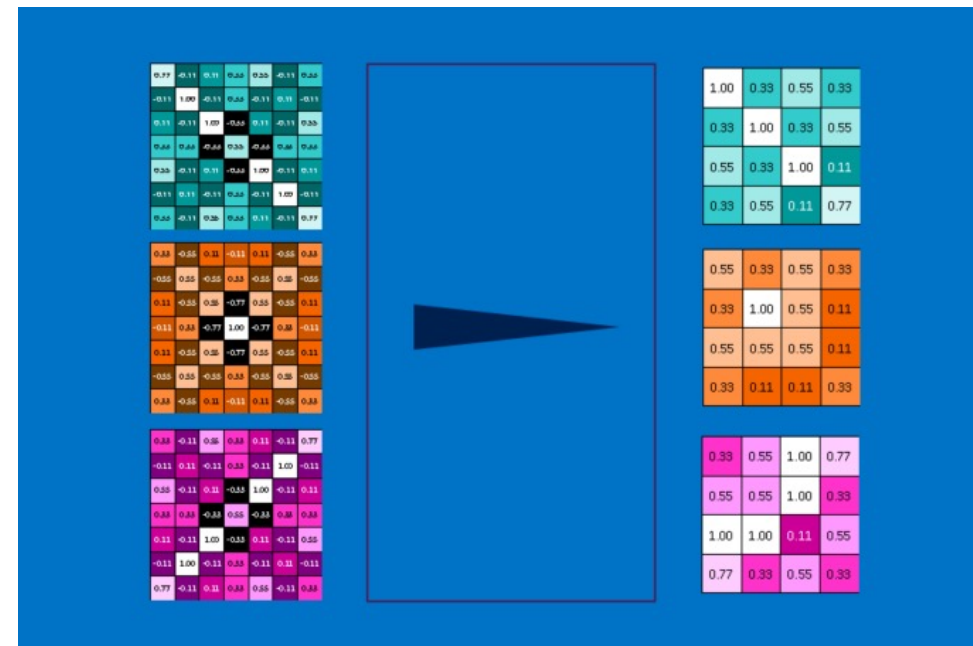

Convolutional Neural Networks

- There are other types of layers besides dense
- A layer used for images is called a **convolutional** layer
 - Apply a set of **filters** to an image
 - Take the output of the filters as the input of a dense neural network
- A filter is just a really small image, 3x3, 4x4, ... pixels
- We apply the filter by going over every 3x3 set of pixels in the original image and seeing how close the original image is to the filter in that region
 - Record the closeness score everywhere
 - This is then the output of the layer



Convolutional Neural Networks

- By applying 3 filters we have tripled the amount of data we have!
- This means in our full network we'll need a lot of neurons on our dense layer
- A common way to fix this is **max pooling**
- We take a small box and cycle it through the output of the filters
- Everywhere the box goes, we just remember the largest number in the box
 - Throw everything else away
 - Don't let the boxes overlap



Convolutional Neural Networks

- For the convolutional step, where did we find the filters?????
- The neural network treats the entries in each pixel of the filters as parameters to be learned
- Back propagation then learns what the filters are
- We still have to pick the size and number of filters

